

SANJAY MURALIDHAR

Data Analytics & Data Science

sanjay.mitblr2023@manipal.learner.edu | +91 7406653204 | [linkedin.com/in/sanjay-muralidhar](https://www.linkedin.com/in/sanjay-muralidhar) | github.com/sanjaym88

SUMMARY

Computer Science & Engineering (Data Science) undergraduate with hands-on experience building data ingestion, parsing, and transformation pipelines that turn unstructured sources into clean, structured, usable data. Comfortable writing and optimizing SQL for extraction and analysis, and building Python-based tools to automate data collection and processing.

TECHNICAL SKILLS

- Data Pipelines & ETL: data ingestion, parsing, and transformation pipelines for unstructured & structured sources
- Databases: MySQL, SQL query writing & optimization, relational database design
- Programming: Python, C
- Data Analysis: Pandas, NumPy, Scikit-learn, feature engineering
- Applied AI: LangChain, Groq API, workflow automation
- Tools: Git, GitHub, Streamlit, Jupyter Notebook

EDUCATION

Manipal Institute of Technology, Bengaluru

Expected 2027

B.Tech, Computer Science & Engineering (Data Science) | CGPA: 7.25/10

Relevant Coursework: Database Management Systems, Data Structures & Algorithms, Machine Learning, Probability & Statistics

EXPERIENCE

Software Development Intern | Tech Mahindra

Jun 2025 – Jul 2025

- Architected a full-stack data pipeline and application (FastAPI backend, REST endpoints) to ingest, parse, and process unstructured documents end-to-end
- Built a document parsing stage to extract structured text, then designed a templated transformation layer that injects source data into structured LLM-based processing
- Integrated streaming LLM inference (Groq LLaMA-3) for real-time output generation, plus a session-based multi-turn interaction layer with server-side state management
- Configured a CORS-enabled API to support a fully self-contained, deployable data application

PROJECTS

Market Sentiment-Based Stock Movement Prediction System | Personal Project

2025

- Built a data pipeline collecting, cleaning, and validating financial news and historical price data from multiple sources using Python and Pandas
- Engineered lagged, price-difference, and rolling-window features to transform raw unstructured text into structured, model-ready records
- Applied statistical methods to improve output quality, raising R^2 from 0.01 with a baseline model to 0.30 with an optimized model
- Identified and resolved data-quality issues, including target leakage, through systematic validation checks
- Built a dashboard to make pipeline output accessible and interpretable for end users